Socio-Informatics 348

Data Visualisation with the Tidyverse
Part 2

Dr Lisa Martin

Department of Information Science
Stellenbosch University

Today's Reading

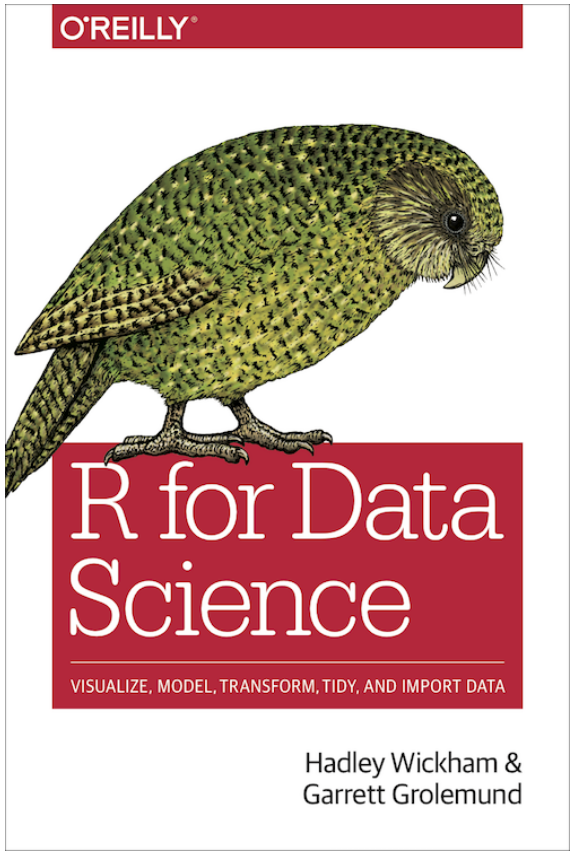

R for Data Science, Wholegame, Visualisation

Side Notes

- **Dark Mode:** Tools ▷ Global Options ▷ Appearance ▷ Editor theme
- **Shortcuts:** Tools ▷ Keyboard Shortcuts Help (Alt+Shift+K)

Where we left off...

- Scatterplot with palmerpenguins dataset
- ggplot layers
- Data, Aesthetics (aes) and Geometry (geom)

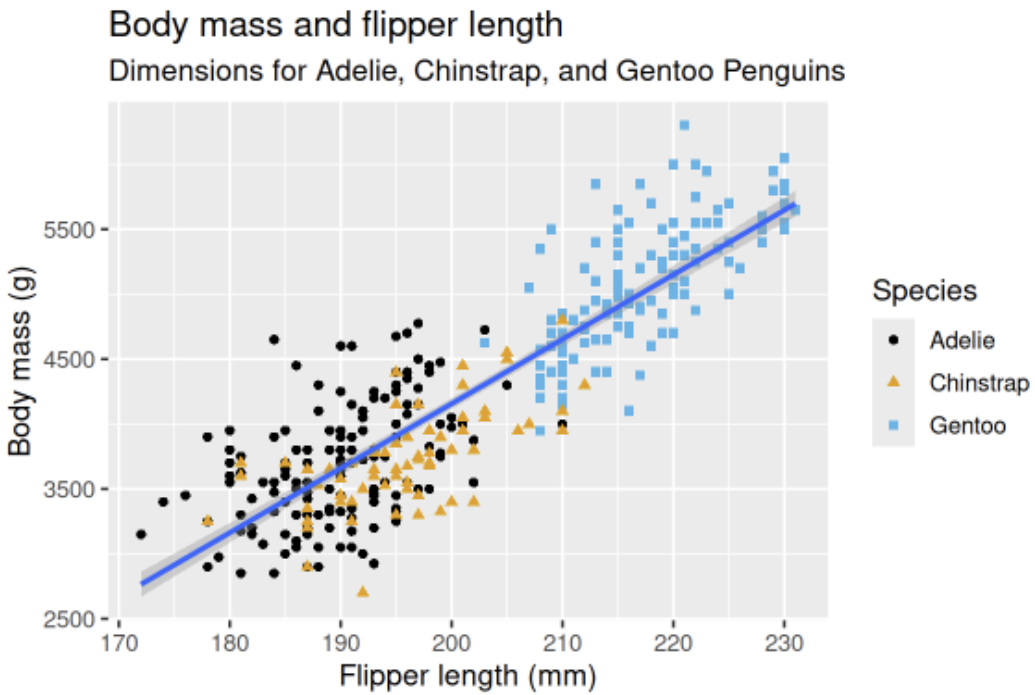

Visualising Distributions

- Understanding the distribution of a variable is important
- Helps to identify patterns, outliers, and the overall shape of the data
- Geom used depends on the type of variable
- Continuous variables: geom_histogram(), geom_density()
- Categorical variables: geom_bar()

Visualising Distributions

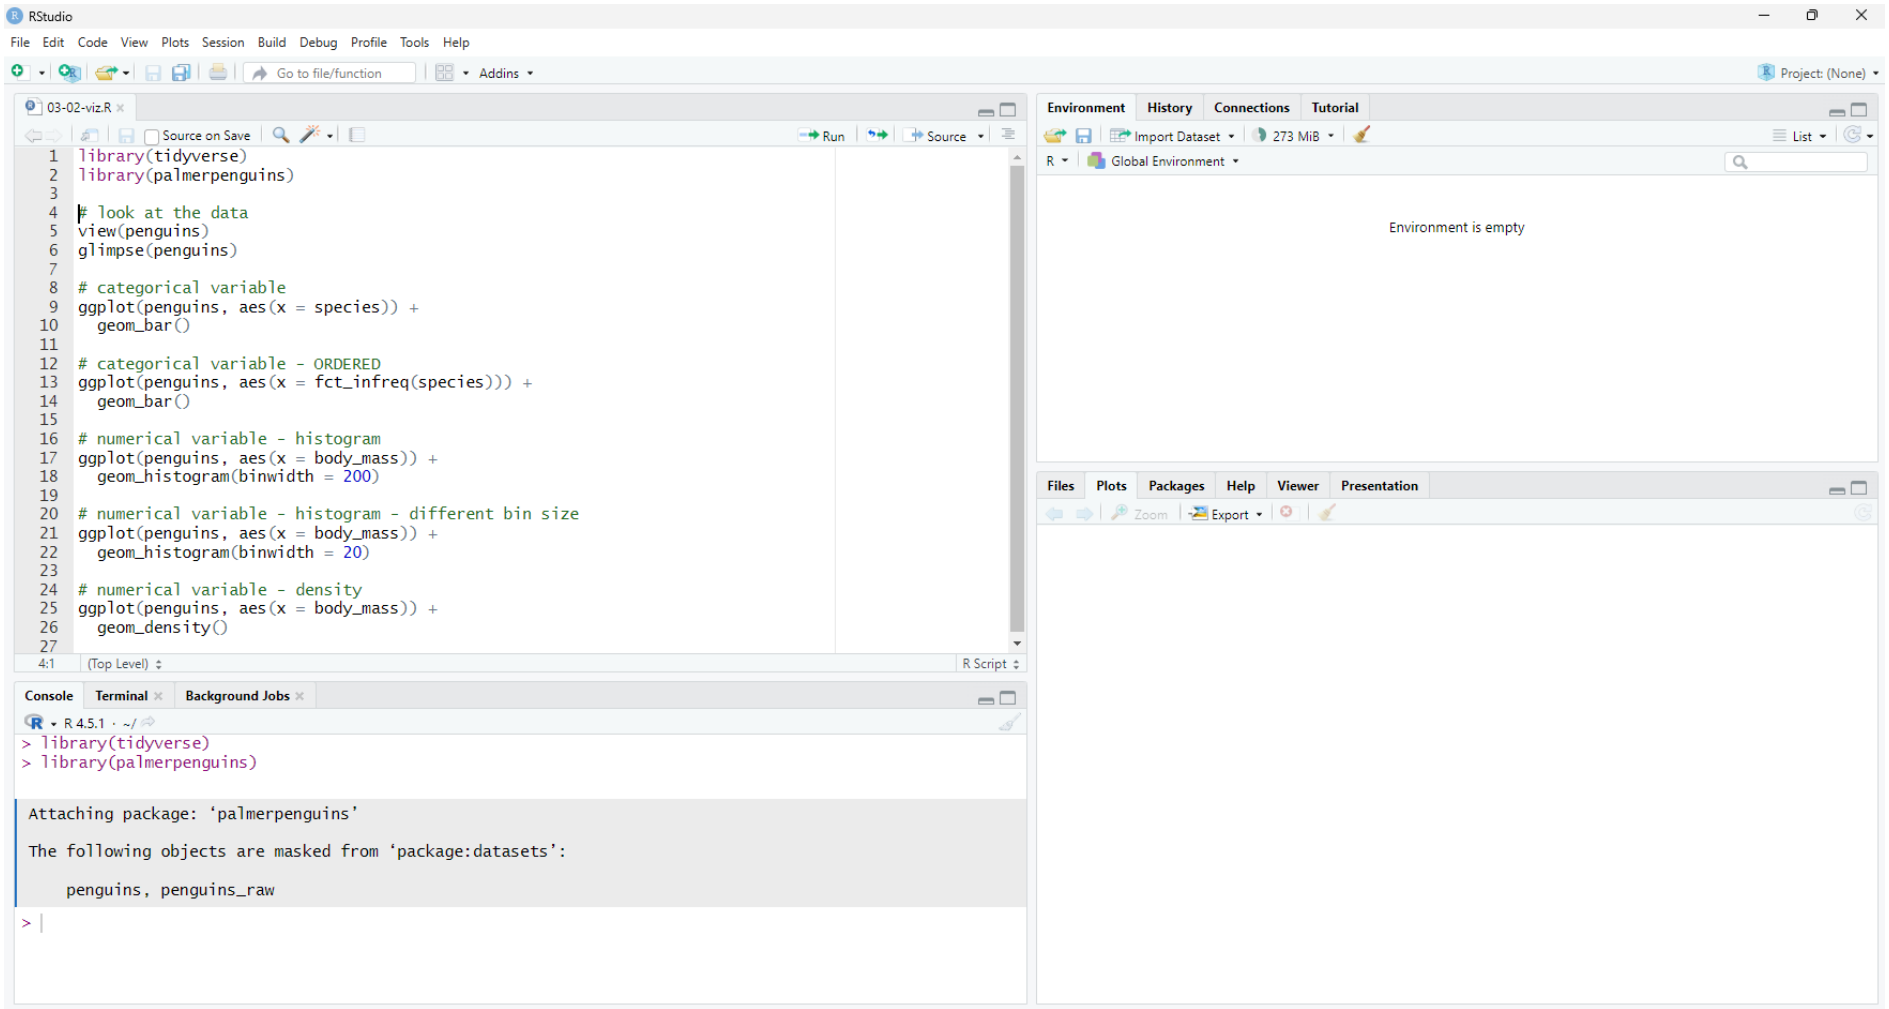

Visualising Distributions

- glimpse()

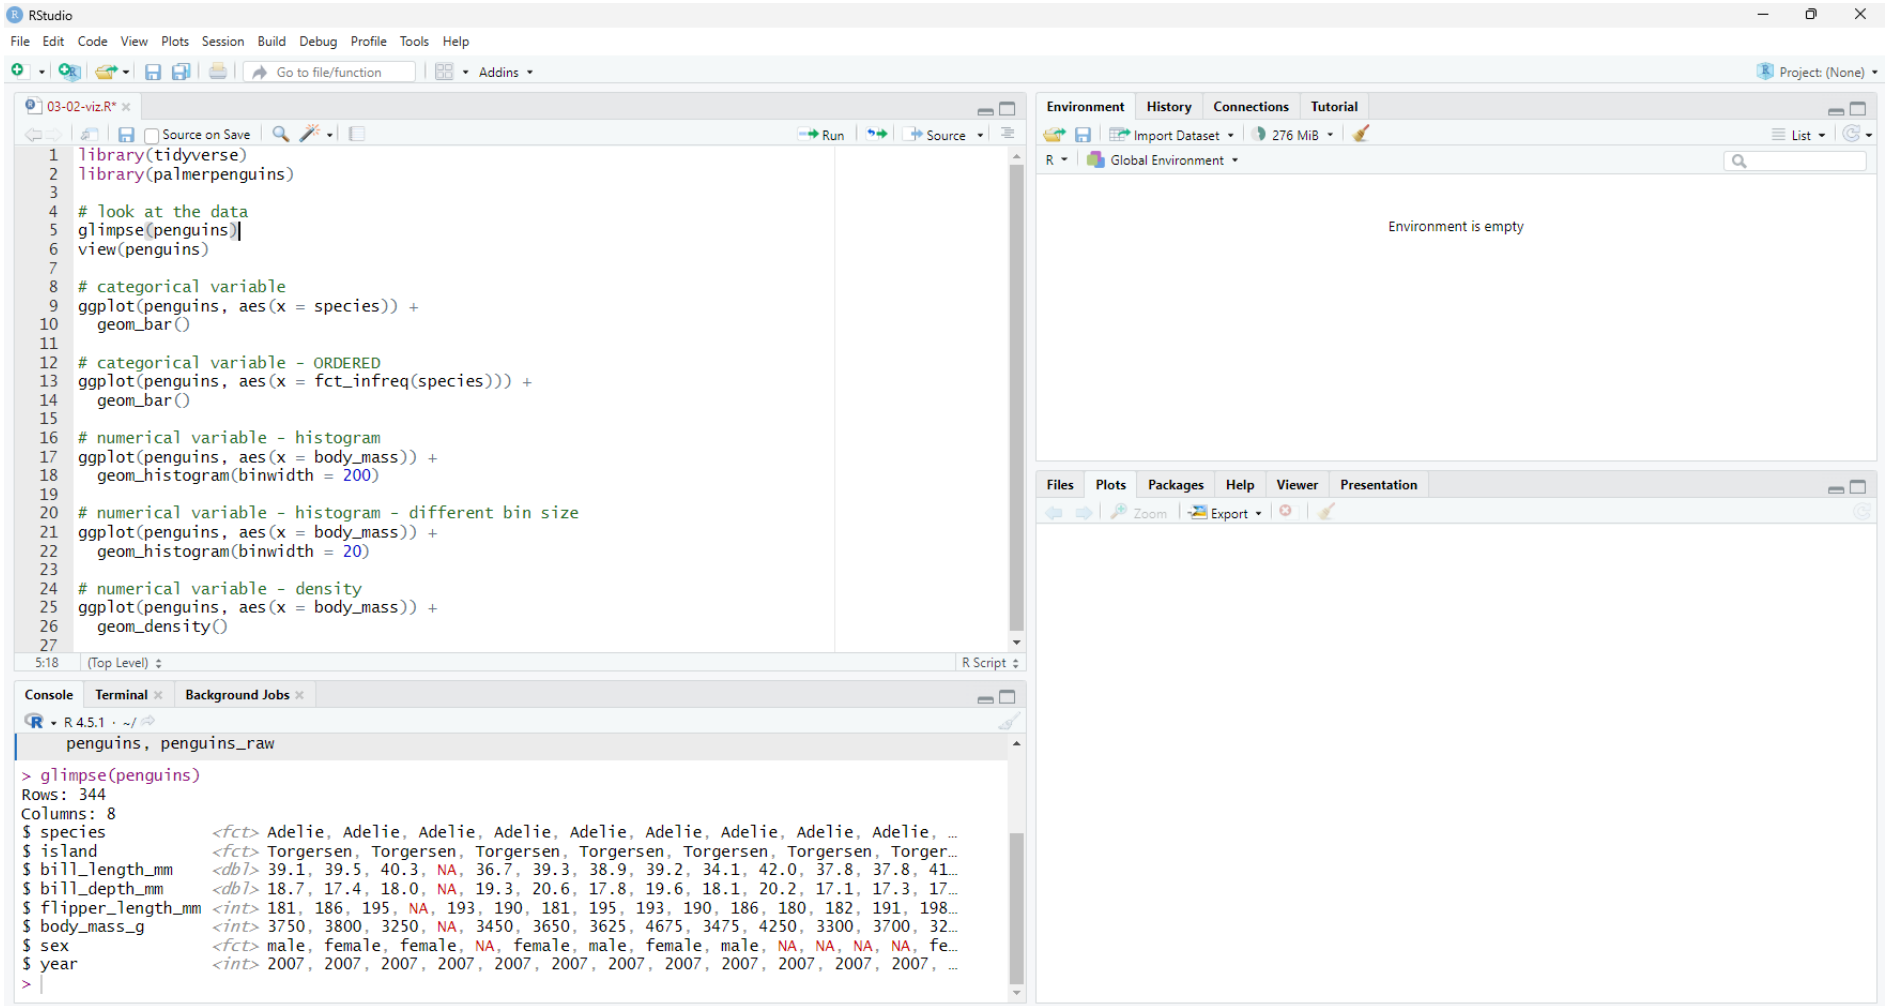

Visualising Distributions

- view()

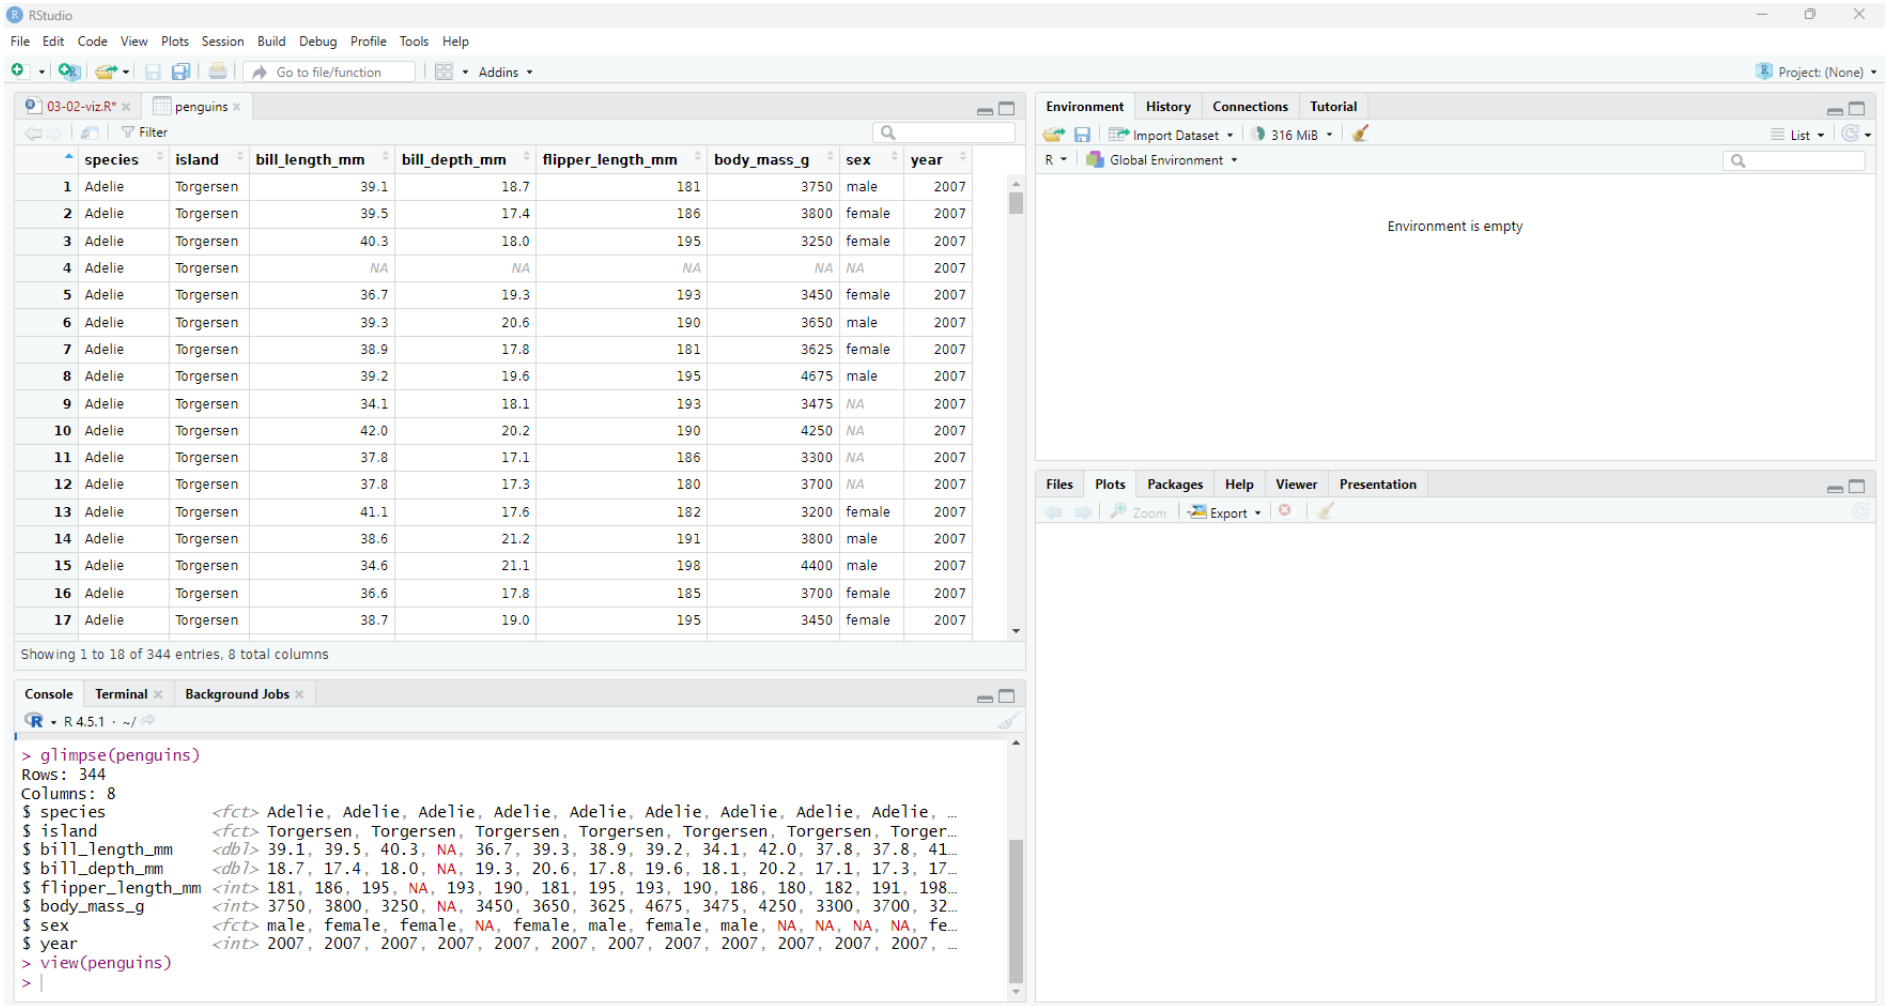

Visualising Distributions

Categorical Variables: geom_bar()

- Can only take one of a small set of values.
- The height of the bars displays how many observations occurred with each x value.
- geom_bar() uses the count of observations by default

Visualising Distributions

Categorical Variables: geom_bar()

```
library(tidyverse)
library(palmerpenguins)

# look at the data
glimpse(penguins)
view(penguins)

# categorical variable
ggplot(penguins, aes(x = species)) +
  geom_bar()

# categorical variable - ORDERED
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass)) +
  geom_density()
```

```
Columns: 8
$ species           <fct> Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, …
$ island            <fct> Torgersen, Torgersen, Torgersen, Torgersen, Torgersen, Torgersen, Torger…
$ bill_length_mm    <dbl> 39.1, 39.5, 40.3, NA, 36.7, 39.3, 38.9, 39.2, 34.1, 42.0, 37.8, 37.8, 41…
$ bill_depth_mm     <dbl> 18.7, 17.4, 18.0, NA, 19.3, 20.6, 17.8, 19.6, 18.1, 20.2, 17.1, 17.3, 17…
$ flipper_length_mm <int> 181, 186, 195, NA, 193, 190, 181, 195, 193, 190, 186, 180, 182, 191, 198…
$ body_mass_g       <int> 3750, 3800, 3250, NA, 3450, 3650, 3625, 4675, 3475, 4250, 3300, 3700, 32…
$ sex               <fct> male, female, female, NA, female, male, female, male, NA, NA, NA, NA, fe…
$ year              <int> 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, …
> view(penguins)
> # categorical variable
> ggplot(penguins, aes(x = species)) +
+   geom_bar()
>
```

Environment is empty

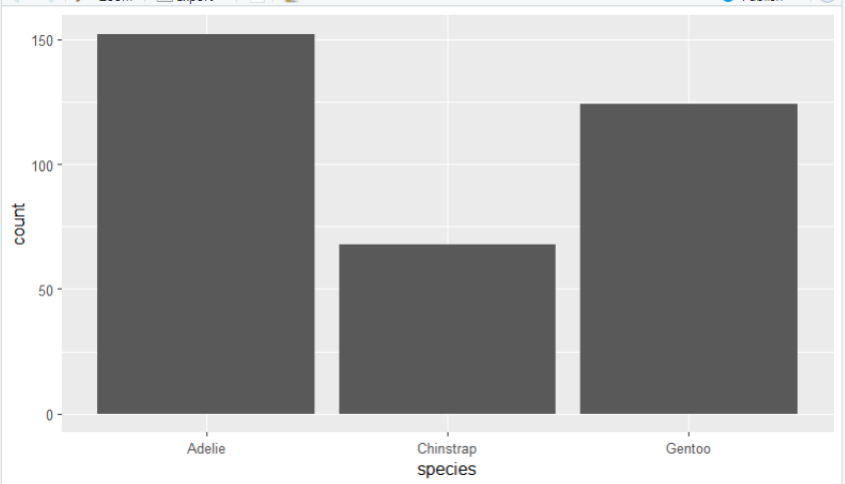

Visualising Distributions

Categorical Variables: geom_bar() - ordered

```
library(tidyverse)
library(palmerpenguins)

# look at the data
glimpse(penguins)
view(penguins)

# categorical variable
ggplot(penguins, aes(x = species)) +
  geom_bar()

# categorical variable - ORDERED
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass)) +
  geom_density()
```

```
$ bill_length_mm    <dbl> 39.1, 39.5, 40.3, NA, 36.7, 39.3, 38.9, 39.2, 34.1, 42.0, 37.8, 37.8, 41…
$ bill_depth_mm     <dbl> 18.7, 17.4, 18.0, NA, 19.3, 20.6, 17.8, 19.6, 18.1, 20.2, 17.1, 17.3, 17…
$ flipper_length_mm <int> 181, 186, 195, NA, 193, 190, 181, 195, 193, 190, 186, 180, 182, 191, 198…
$ body_mass_g       <int> 3750, 3800, 3250, NA, 3450, 3650, 3625, 4675, 3475, 4250, 3300, 3700, 32…
$ sex               <fct> male, female, female, NA, female, male, female, male, NA, NA, NA, NA, fe…
$ year              <int> 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, …
> view(penguins)
> # categorical variable
> ggplot(penguins, aes(x = species)) +
+   geom_bar()
> # categorical variable - ORDERED
> ggplot(penguins, aes(x = fct_infreq(species))) +
+   geom_bar()
>
```

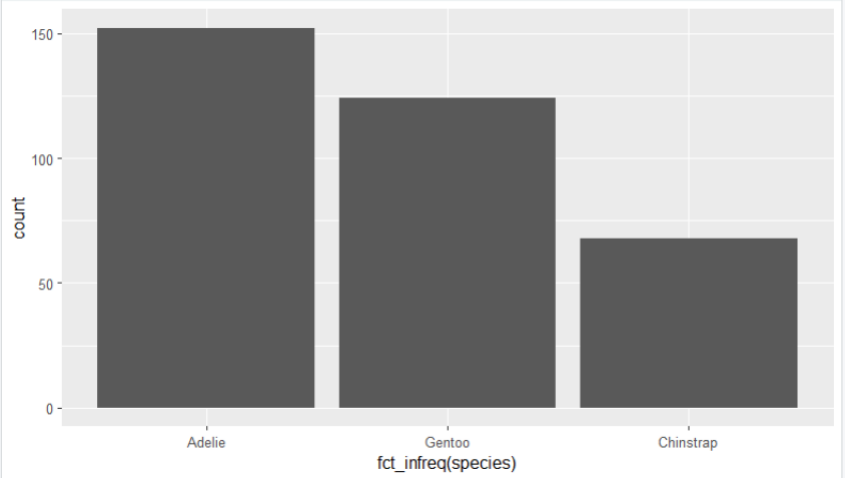

Visualising Distributions

Numerical Variables: geom_histogram()

```
library(tidyverse)
library(palmerpenguins)

# look at the data
glimpse(penguins)
view(penguins)

# categorical variable
ggplot(penguins, aes(x = species)) +
  geom_bar()

# categorical variable - ORDERED
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()
```

```
> # categorical variable - ORDERED
> ggplot(penguins, aes(x = fct_infreq(species))) +
+   geom_bar()
> # numerical variable - histogram
> ggplot(penguins, aes(x = body_mass_g)) +
+   geom_histogram(binwidth = 200)
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_bin()`).

> sum(is.na(penguins$body_mass_g))
[1] 2
>
```

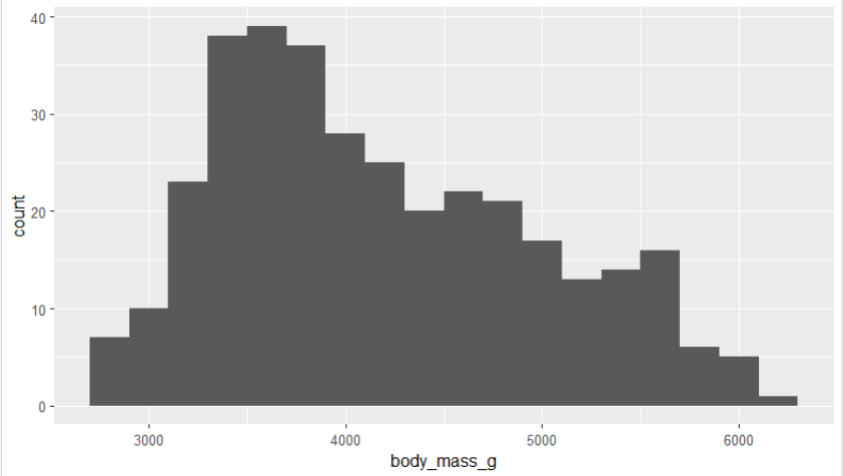

Note: Warning about NAs

Visualising Distributions

Numerical Variables: geom_histogram() - bin size

```
library(tidyverse)
library(palmerpenguins)

# look at the data
glimpse(penguins)
view(penguins)

# categorical variable
ggplot(penguins, aes(x = species)) +
  geom_bar()

# categorical variable - ORDERED
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()
```

```
Removed 2 rows containing non-finite outside the scale range (`stat_bin()`).

> sum(is.na(penguins$body_mass_g))
[1] 2
> # numerical variable - histogram - different bin size
> ggplot(penguins, aes(x = body_mass_g)) +
+   geom_histogram(binwidth = 20)

Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_bin()`).
```

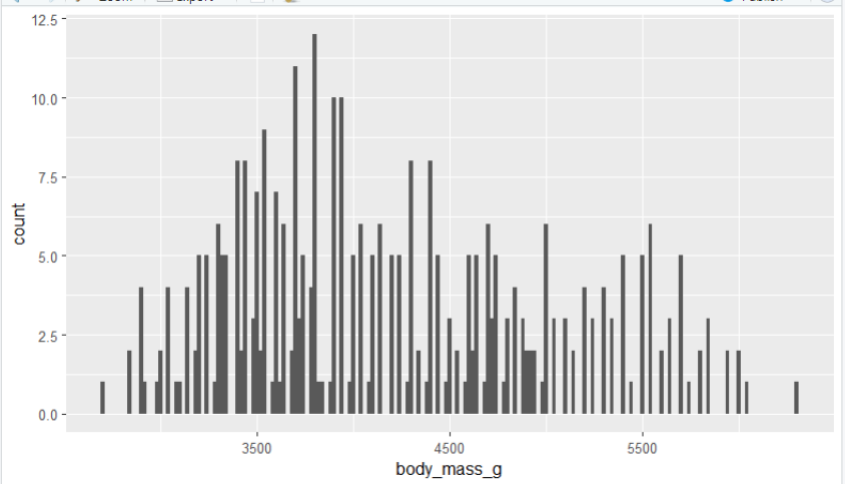

Visualising Distributions

Numerical Variables: geom_density()

- Smoothed version of a histogram
- Good for continuous data

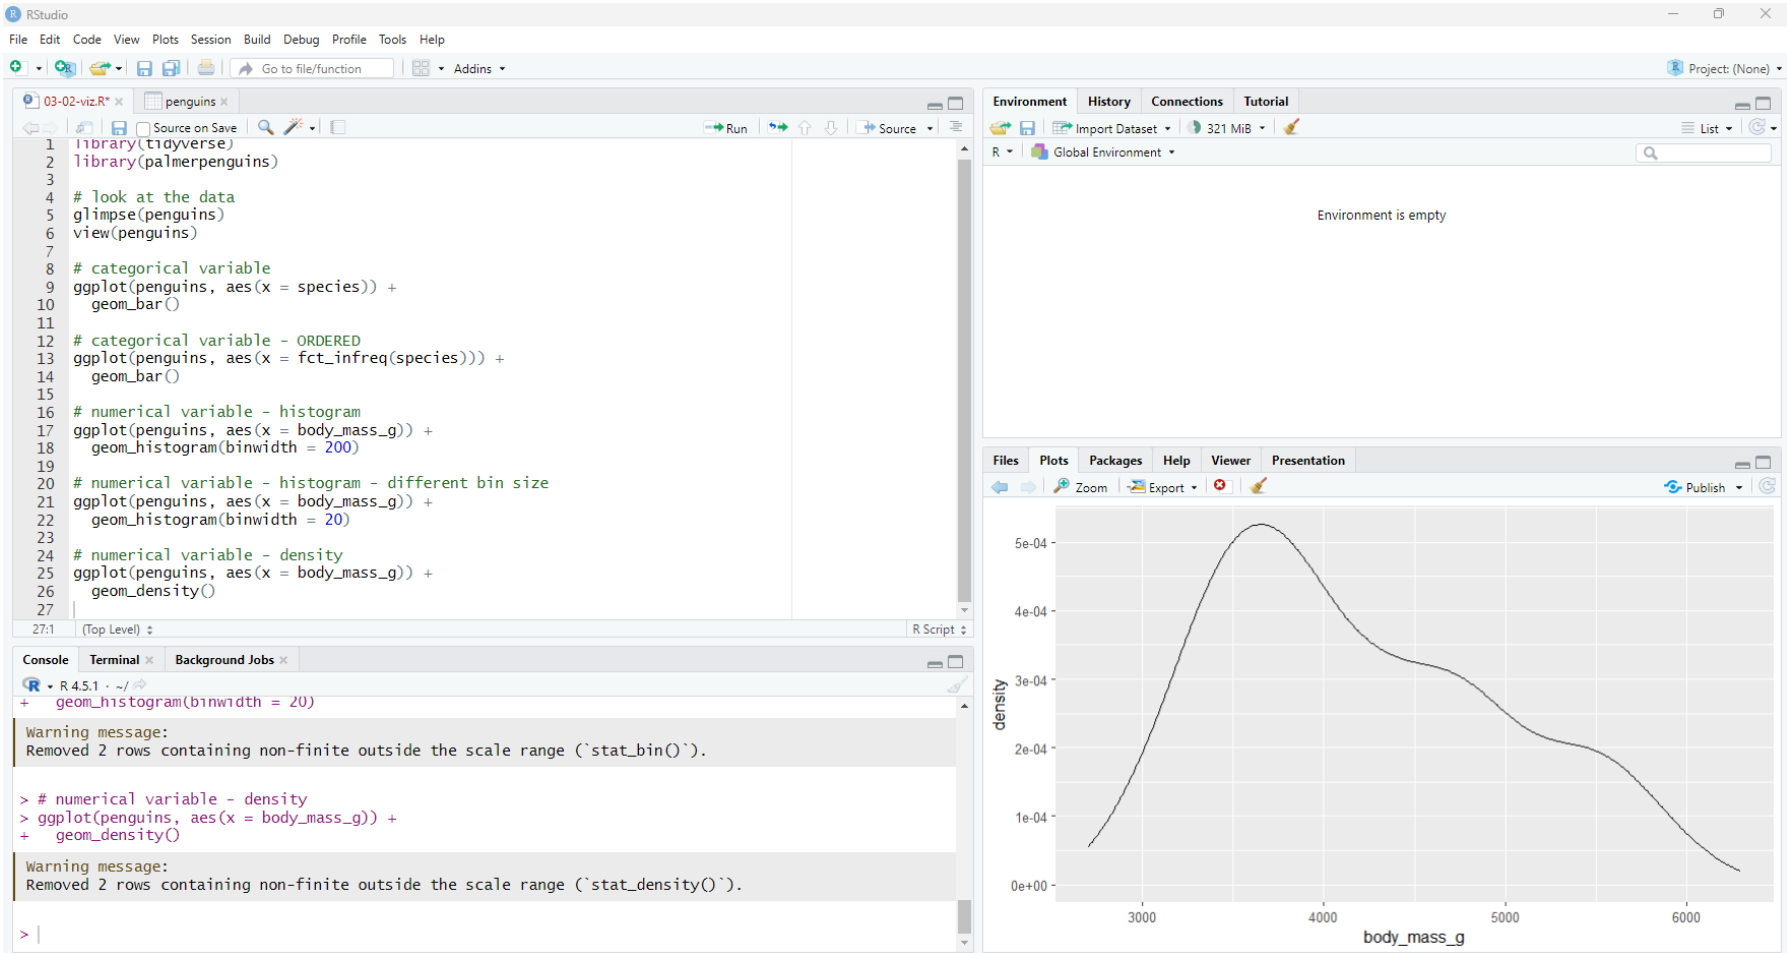

Visualising Relationships

Numerical and Categorical: geom_boxplot()

- Visualise the distribution of a continuous variable across categories
- Displays the median, quartiles, and potential outliers
- Spread of the distribution - symmetric or skewed to one side

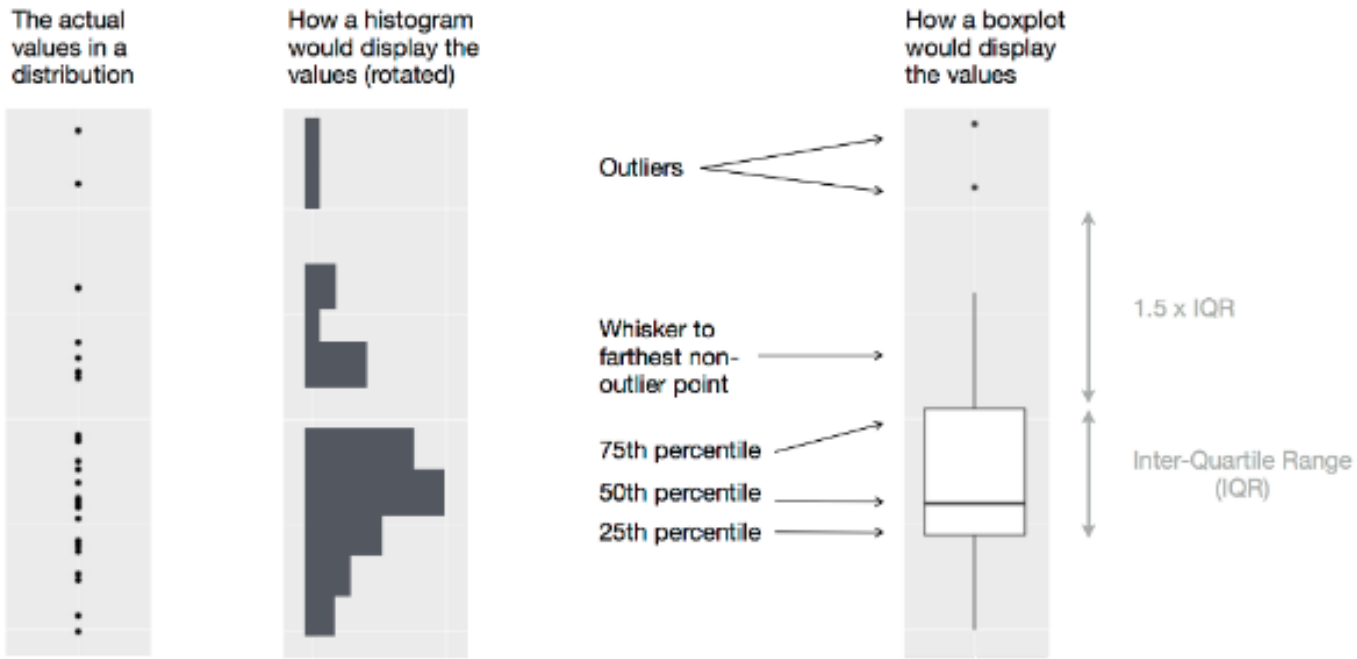

R4DS, Figure 1.1

Visualising Relationships

Numerical and Categorical: geom_boxplot()

```
glimpse(penguins)
view(penguins)

# categorical variable
ggplot(penguins, aes(x = species)) +
  geom_bar()

# categorical variable - ORDERED
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()

# boxplot
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()
```

```
+   geom_density()
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_density()`).

> # boxplot
> ggplot(penguins, aes(x = species, y = body_mass_g)) +
+   geom_boxplot()
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_boxplot()`).

>
```

Environment is empty

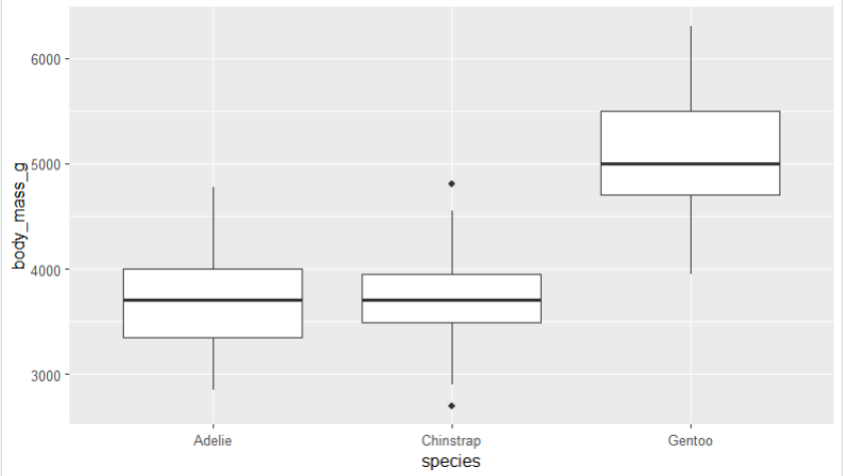

Visualising Relationships

Numerical and Categorical: geom_density()

```
# categorical variable - ORDERED
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()

# boxplot
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()

# density
ggplot(penguins, aes(x = body_mass_g, color = species)) +
  geom_density(linewidth = 0.75)
```

```
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_boxplot()`).

> # density
> ggplot(penguins, aes(x = body_mass_g, color = species)) +
+   geom_density(linewidth = 0.75)

Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_density()`).
```

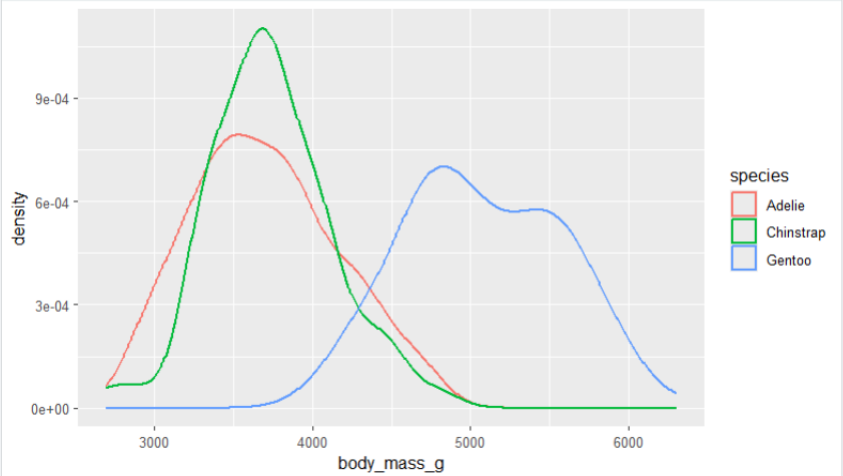

Note: linewidth to set the thickness of the line

Visualising Relationships

Numerical and Categorical: geom_density()

```
ggplot(penguins, aes(x = fct_infreq(species))) +
  geom_bar()

# numerical variable - histogram
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 200)

# numerical variable - histogram - different bin size
ggplot(penguins, aes(x = body_mass_g)) +
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()

# boxplot
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()

# density
ggplot(penguins, aes(x = body_mass_g, color = species)) +
  geom_density(linewidth = 0.75)

# use linewidth and fill
ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
  geom_density(alpha = 0.5)
```

```
+   geom_density(linewidth = 0.75)
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_density()`).

> # use linewidth and fill
> ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
+   geom_density(alpha = 0.5)
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_density()`).

>
```

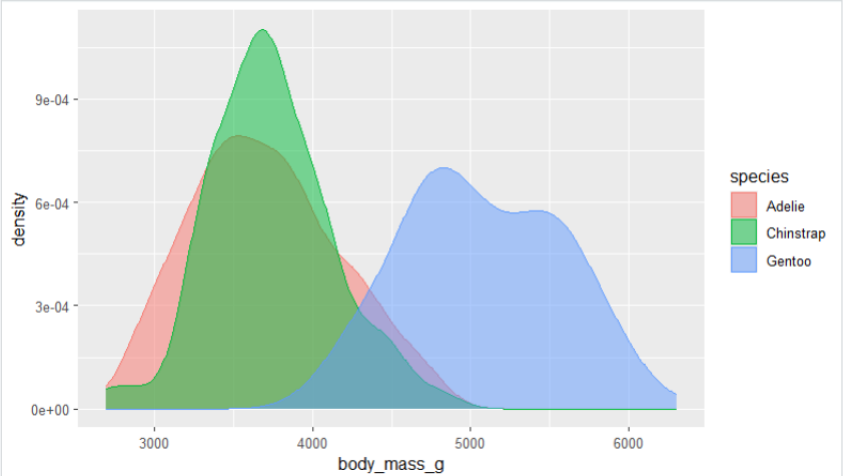

Note: linewidth and fill to set the thickness and colour of the area under the curve
alpha to set the transparency of the fill

Visualising Relationships

Two Categorical: geom_bar()

```
# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()

# RELATIONSHIP
# numerical variable - boxplot
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g, color = species)) +
  geom_density(linewidth = 0.75)

#  numerical variable - density - use linewidth and fill
ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
  geom_density(alpha = 0.5)

# 2 categorical variables
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar()

# 2 categorical variables - percentages
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar(position = "fill")
```

```
> # use linewidth and fill
> ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
+   geom_density(alpha = 0.5)
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_density()`).

> # 2 categorical variables
> ggplot(penguins, aes(x = island, fill = species)) +
+   geom_bar()
>
```

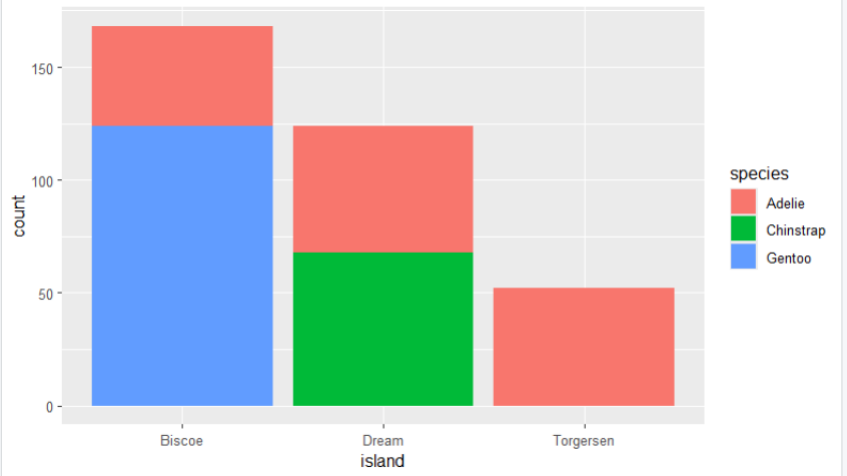

Visualising Relationships

Two Categorical: geom_bar()

```
  geom_histogram(binwidth = 20)

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g)) +
  geom_density()

# RELATIONSHIP
# numerical variable - boxplot
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g, color = species)) +
  geom_density(linewidth = 0.75)

#  numerical variable - density - use linewidth and fill
ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
  geom_density(alpha = 0.5)

# 2 categorical variables
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar()

# 2 categorical variables - percentages
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar(position = "fill")
```

```
> ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
+   geom_density(alpha = 0.5)
Warning message:
Removed 2 rows containing non-finite outside the scale range (`stat_density()`).

> # 2 categorical variables
> ggplot(penguins, aes(x = island, fill = species)) +
+   geom_bar()
> # 2 categorical variables - percentages
> ggplot(penguins, aes(x = island, fill = species)) +
+   geom_bar(position = "fill")
>
```

Environment is empty

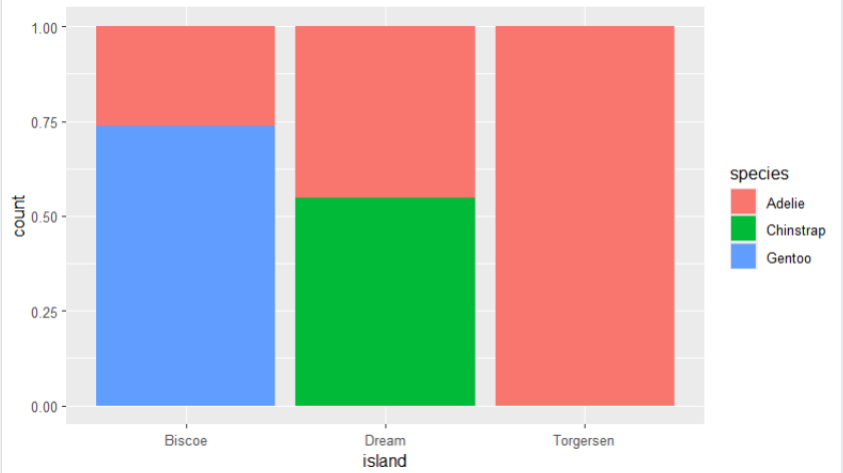

Note: Use of position = "fill" to show proportions instead of counts

Visualising Relationships

Two Numerical: geom_point()

```
  geom_density()

# RELATIONSHIP
# numerical variable - boxplot
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g, color = species)) +
  geom_density(linewidth = 0.75)

#  numerical variable - density - use linewidth and fill
ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
  geom_density(alpha = 0.5)

# 2 categorical variables
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar()

# 2 categorical variables - percentages
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar(position = "fill")

# 2 numerical variables
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point()
```

```
> ggplot(penguins, aes(x = island, fill = species)) +
+   geom_bar()
> # 2 categorical variables - percentages
> ggplot(penguins, aes(x = island, fill = species)) +
+   geom_bar(position = "fill")
> # 2 numerical variables
> ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
+   geom_point()
Warning message:
Removed 2 rows containing missing values or values outside the scale range (`geom_point()`).
>
```

Environment is empty

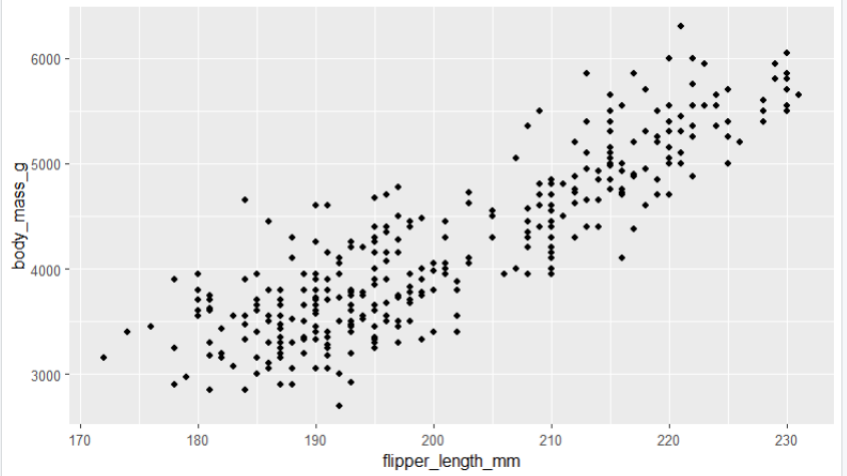

Visualising Relationships

Two or more variables: Additional aesthetics

```
ggplot(penguins, aes(x = species, y = body_mass_g)) +
  geom_boxplot()

# numerical variable - density
ggplot(penguins, aes(x = body_mass_g, color = species)) +
  geom_density(linewidth = 0.75)

#  numerical variable - density - use linewidth and fill
ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
  geom_density(alpha = 0.5)

# 2 categorical variables
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar()

# 2 categorical variables - percentages
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar(position = "fill")

# 2 numerical variables
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point()

# 3 or more variables
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point(aes(color = species, shape = island))
```

```
+   geom_point()

Warning message:
Removed 2 rows containing missing values or values outside the scale range (`geom_point()`).

> # 3 or more variables
> ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
+   geom_point(aes(color = species, shape = island))

Warning message:
Removed 2 rows containing missing values or values outside the scale range (`geom_point()`).

>
```

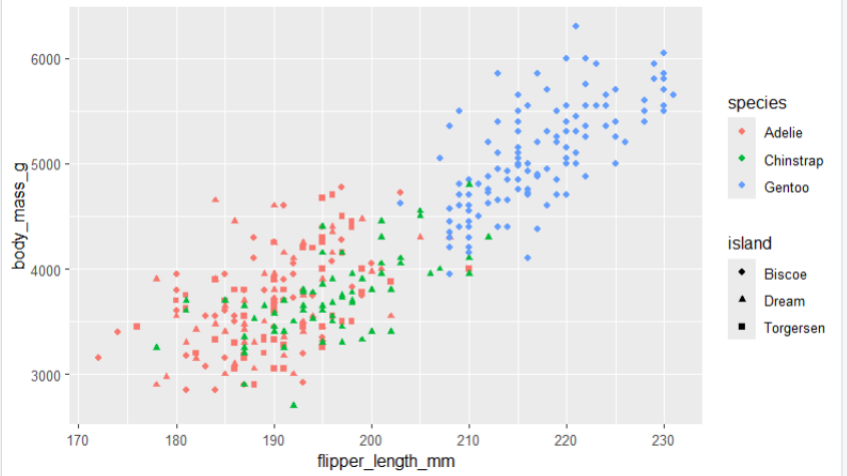

Visualising Relationships

Two or more variables: Facets

```
  geom_density(linewidth = 0.75)

#  numerical variable - density - use linewidth and fill
ggplot(penguins, aes(x = body_mass_g, color = species, fill = species)) +
  geom_density(alpha = 0.5)

# 2 categorical variables
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar()

# 2 categorical variables - percentages
ggplot(penguins, aes(x = island, fill = species)) +
  geom_bar(position = "fill")

# 2 numerical variables
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point()

# 3 or more variables - aesthetics
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point(aes(color = species, shape = island))

# 3 or more variables - facets
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point(aes(color = species, shape = species)) +
  facet_wrap(~island)
```

```
Warning message:
Removed 2 rows containing missing values or values outside the scale range (`geom_point()`).

> # 3 or more variables - facets
> ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
+   geom_point(aes(color = species, shape = species)) +
+   facet_wrap(~island)

Warning message:
Removed 2 rows containing missing values or values outside the scale range (`geom_point()`).
```

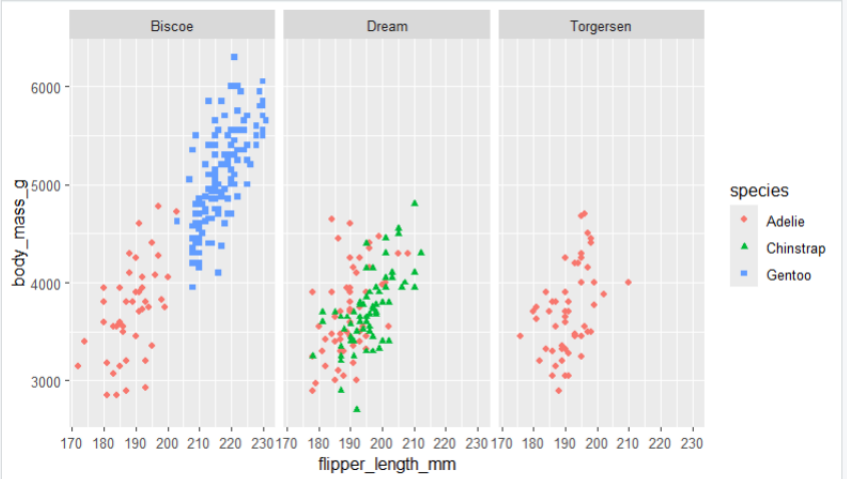

Save your plots

ggsave()

```
ggplot(penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point()
ggsave(filename = "penguin-plot.png")
```