

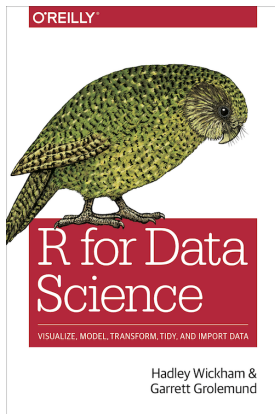
Socio-Informatics 348

Data Transformation Regex

Dr Lisa Martin

Department of Information Science
Stellenbosch University

Today's Reading



R for Data Science, Chapter 15

Why Regular Expressions?

- Regular expressions (regex/regexp): concise, powerful language for string patterns.
- Use cases: extract, detect, count, replace text.

Pattern Basics

Use `str_view()` to do basic pattern matching.

- **Literal characters:** match themselves exactly (e.g., `berry`).

```
str_view(fruit, "berry")  
#> [6] | bil<berry>  
#> [7] | black<berry>  
#> [10] | blue<berry>  
#> [11] | boysen<berry>  
#> [19] | cloud<berry>  
#> [21] | cran<berry>  
#> ... and 8 more
```

- **Metacharacters:** Punctuation and special meanings
- **Quantifiers:** How many times can a pattern match?

Pattern Basics

- **Literal characters:** match themselves exactly (e.g., berry).
- **Metacharacters:** Punctuation and special characters
. + * [] ? etc.
- **Dot (.):** any single character.

```
str_view(c("a", "ab", "ae", "bd", "ea", "eab"), "a.")  
#> [2] | <ab>  
#> [3] | <ae>  
#> [6] | e<ab>
```

- **Quantifiers:** How many times can a pattern match?

Pattern Basics

- **Literal characters:** match themselves exactly (e.g., berry).
- **Metacharacters:** Punctuation and special characters
. + * [] ? etc.
- **Dot (.):** any single character.

```
str_view(fruit, "a...e")  
#> [1] | <apple>  
#> [7] | bl<ackbe>rry  
#> [48] | mand<arine>  
#> [51] | nect<arine>  
#> [62] | pine<apple>  
#> [64] | pomegr<anate>  
#> ... and 2 more
```

- **Quantifiers:** How many times can a pattern match?

Pattern Basics

- **Quantifiers:** How many times can a pattern match?
? optional, + one or more, * zero or more.

```
# ab? matches an "a", optionally followed by a "b".
```

```
str_view(c("a", "ab", "abb"), "ab?")
```

```
#> [1] | <a>
```

```
#> [2] | <ab>
```

```
#> [3] | <ab>b
```

```
# ab+ matches an "a", followed by at least one "b".
```

```
str_view(c("a", "ab", "abb"), "ab+")
```

```
#> [2] | <ab>
```

```
#> [3] | <abb>
```

```
# ab* matches an "a", followed by any number of "b"s.
```

```
str_view(c("a", "ab", "abb"), "ab*")
```

```
#> [1] | <a>
```

```
#> [2] | <ab>
```

```
#> [3] | <abb>
```

Character Sets and Alternation

- Character classes are defined by square brackets: `[abcd]`.
- Find the words containing an "x" surrounded by vowels, or a "y" surrounded by consonants:

```
str_view(words, "[aeiou]x[aeiou]")
#> [284] | <exa>ct
#> [285] | <exa>mple
#> [288] | <exe>rcise
#> [289] | <exi>st
str_view(words, "[^aeiou]y[^aeiou]")
#> [836] | <sys>tem
#> [901] | <typ>e
```

- Note: invert the match with `^` inside brackets: `[^aeiou]`.

Character Sets and Alternation

- Alternation is defined by vertical bar: `a|b|c`.
- Look for fruits containing "apple", "melon", or "nut", or a repeated vowel:

```
str_view(fruit, "apple|melon|nut")
```

```
#> [1] | <apple>  
#> [13] | canary <melon>  
#> [20] | coco<nut>  
#> [52] | <nut>  
#> [62] | pine<apple>  
#> [72] | rock <melon>  
#> ... and 1 more
```

```
str_view(fruit, "aa|ee|ii|oo|uu")
```

```
#> [9] | bl<oo>d orange  
#> [33] | g<oo>seberry  
#> [47] | lych<ee>  
#> [66] | purple mangost<ee>n
```

Key Regex Functions

- `str_detect()` – logical vector.
- `str_subset()` – return matches.
- `str_which()` – positions.
- `str_count()` – number of matches.
- `str_replace()`, `str_replace_all()`.
- `separate_wider_regex()` – extract to columns.

str_detect()

- `str_detect(string, pattern)` returns TRUE/FALSE if pattern is found in string.

```
str_detect(c("a", "b", "c"), "[aeiou]")  
#> [1] TRUE FALSE FALSE
```


- Useful for filtering with `filter()`.

```
babynames |>  
  filter(str_detect(name, "x")) |>  
  count(name, wt = n, sort = TRUE)
```

str_detect()

- useful with `sum()` and `mean()`.

```
babynames |>  
  group_by(year) |>  
  summarize(prop_x = mean(str_detect(name, "x"))) |>  
  ggplot(aes(x = year, y = prop_x)) +  
  geom_line()
```



str_count()

- `str_count(string, pattern)` counts occurrences of pattern in string.

```
x <- c("apple", "banana", "pear")  
str_count(x, "p")  
#> [1] 2 0 1
```

- Note that patterns cannot overlap.

```
str_count("abababa", "aba")  
#> [1] 2  
str_view("abababa", "aba")  
#> [1] | <aba>b<aba>
```

str_count()

- Use with `mutate()` to create new variables.

```
babynames |>
  count(name) |>
  mutate(
    vowels = str_count(name, "[aeiou]"),
    consonants = str_count(name, "[^aeiou]")
  )
#> # A tibble: 97,310 × 4
#>   name          n vowels consonants
#>   <chr>      <int>  <int>      <int>
#> 1 Aaban         10     2          3
#> 2 Aabha          5     2          3
#> 3 Aabid          2     2          3
#> 4 Aabir          1     2          3
#> 5 Aabriella      5     4          5
#> 6 Aada           1     2          2
#> # i 97,304 more rows
```

- What do you notice about the number of vowels?

str_count()

Solve the issue:

- Include uppercase vowels to character class
`str_count(name, "[aeiouAEIOU]").`
- Tell the function to ignore case
`str_count(name, regex("[aeiou]", ignore_case = TRUE)).`
- Convert names to lower case before counting
`str_count(str_to_lower(name), "[aeiou]").`

str_replace() and str_remove()

- Replace first match:
`str_replace(string, pattern, replacement)`
- Replace all matches:
`str_replace_all(string, pattern, replacement)`

```
x <- c("apple", "pear", "banana")  
str_replace_all(x, "[aeiou]", "-")  
#> [1] "-pp1-" "p--r"  "b-n-n-"
```

- Remove matches: set replacement to empty string "" OR use `str_remove()` and `str_remove_all()`.

```
x <- c("apple", "pear", "banana")  
str_remove_all(x, "[aeiou]")  
#> [1] "pp1" "pr"  "bnn"
```

- Very useful for data cleaning with `mutate()`.

Extract variables with `separate_wider_regex()`

- Similar to `separate_wider_position` and `separate_wider_delim`.

```
df <- tribble(  
  ~str,  
  "<Sheryl>-F_34",  
  "<Kisha>-F_45",  
  "<Brandon>-N_33",  
  "<Sharon>-F_38",  
  "<Penny>-F_58",  
  "<Justin>-M_41",  
  "<Patricia>-F_84",  
)
```

Extract variables with `separate_wider_regex()`

- Similar to `separate_wider_position` and `separate_wider_delim`.

```
df |>
  separate_wider_regex(
    str,
    patterns = c(
      "<",
      name = "[A-Za-z]+",
      ">-",
      gender = ".",
      "_",
      age = "[0-9]+"
    )
  )
#> # A tibble: 7 × 3
#>   name    gender age
#>   <chr>   <chr> <chr>
#> 1 Sheryl F      34
#> 2 Kisha  F      45
#> 3 Brandon N      33
#> 4 Sharon F      38
#> 5 Penny  F      58
```

Escaping

- Escape metacharacters with backslash – can get tricky!

```
x <- "a\\b"  
str_view(x)  
#> [1] | a\b  
str_view(x, "\\\\" )  
#> [1] | a<\>b
```

- Alternative approaches – raw strings and character classes.

```
str_view(x, r"{\\}")  
#> [1] | a<\>b
```

```
str_view(c("abc", "a.c", "a*c", "a c"), "a[.]c")  
#> [2] | <a.c>  
str_view(c("abc", "a.c", "a*c", "a c"), ".[*]c")  
#> [3] | <a*c>
```

Anchors and Word Boundaries

- Anchors: `^` (start), `$` (end).

```
str_view(fruit, "^a")
```

```
#> [1] | <a>pple
```

```
#> [2] | <a>pricot
```

```
#> [3] | <a>vocado
```

```
str_view(fruit, "a$")
```

```
#> [4] | banan<a>
```

```
#> [15] | cherimoy<a>
```

```
#> [30] | feijo<a>
```

```
#> [36] | guav<a>
```

```
#> [56] | papay<a>
```

```
#> [74] | satsum<a>
```

- Full-string match: `^pattern$`.

Anchors and Word Boundaries

- Word boundary:

```
x <- c("summary(x)", "summarize(df)", "rowsum(x)", "sum(x)")
str_view(x, "sum")
#> [1] | <sum>mary(x)
#> [2] | <sum>marize(df)
#> [3] | row<sum>(x)
#> [4] | <sum>(x)
str_view(x, "\\bsum\\b")
#> [4] | <sum>(x)
```

More on character sets/classes

- - (range)
- ^ (negation).
- \ (escape).

```
x <- "abcd ABCD 12345 -!@#%."
str_view(x, "[abc]+")
#> [1] | <abc>d ABCD 12345 -!@#%.
str_view(x, "[a-z]+")
#> [1] | <abcd> ABCD 12345 -!@#%.
str_view(x, "[^a-z0-9]+")
#> [1] | abcd< ABCD >12345< -!@#%.>
```

```
# You need an escape to match characters that are otherwise
# special inside of []
```

```
str_view("a-b-c", "[a-c]")
#> [1] | <a>-<b>-<c>
str_view("a-b-c", "[a\\-c]")
#> [1] | <a><->b<->c<
```

More on character sets/classes

- Shorthand: `\d` digit, `\s` whitespace, `\w` word char.
- Uppercase negates: `\D`, `\S`, `\W`.

```
x <- "abcd ABCD 12345 -!@#%."
str_view(x, "\\d+")
#> [1] | abcd ABCD <12345> -!@#%.
str_view(x, "\\D+")
#> [1] | <abcd ABCD >12345< -!@#%.>
str_view(x, "\\s+")
#> [1] | abcd< >ABCD< >12345< >-!@#%.
str_view(x, "\\S+")
#> [1] | <abcd> <ABCD> <12345> <-!@#%.>
str_view(x, "\\w+")
#> [1] | <abcd> <ABCD> <12345> -!@#%.
str_view(x, "\\W+")
#> [1] | abcd< >ABCD< >12345< -!@#%.>
```

Quantifiers

- Remember: `?`, `+`, `*`.
- Exact counts: `{n}`. E.g. `\d{3}` for exactly 3 digits.
- At least n: `{n,}`. E.g. `\d{3,}` for 3 or more digits.
- Range: `{n,m}`. E.g. `\d{3,5}` for between 3 and 5 digits.

Operator Precedence

- Precedence: quantifiers $>$ concatenation $>$ alternation.
- Example: $ab^+ = a(b^+)$.
- Example: $^a|b\$ = (^a)|(b\$)$.
- Use parentheses for grouping.

Grouping and Capturing

- Parentheses capture submatches/groupings.
- Backreference refers to grouping: \1, \2, ...

```
str_view(fruit, "(..)\1")
```

```
#> [4] | b<anan>a  
#> [20] | <coco>nut  
#> [22] | <cucu>mber  
#> [41] | <juju>be  
#> [56] | <papa>ya  
#> [73] | s<alal> berry
```

```
str_view(words, "^(..).*\1$")
```

```
#> [152] | <church>  
#> [217] | <decide>  
#> [617] | <photograph>  
#> [699] | <require>  
#> [739] | <sense>
```

Pattern Control

- Flags: `ignore_case=TRUE`, `dotall=TRUE`, `multiline=TRUE`, `comments=TRUE`.
- Use `regex()` to set flags **around a pattern**.
- See more in `?regex`.

```
bananas <- c("banana", "Banana", "BANANA")
str_view(bananas, "banana")
#> [1] | <banana>
str_view(bananas, regex("banana", ignore_case = TRUE))
#> [1] | <banana>
#> [2] | <Banana>
#> [3] | <BANANA>
```

Pattern Control

- `fixed()` for literal matches.

```
str_view(c("", "a", "."), fixed("."))  
#> [3] | <.>
```

```
str_view("x X", "X")  
#> [1] | x <X>  
str_view("x X", fixed("X", ignore_case = TRUE))  
#> [1] | <x> <X>
```