


Socio-Informatics 348

Supervised Machine Learning Modelling Overview

Dr Lisa Martin

Department of Information Science
Stellenbosch University

Today's Reading



Computational Analysis of Communication, Chapter 8

Statistical Modeling and Prediction

- Supervised Machine Learning (SML) – a form of machine learning
- SML shares many methods with classical statistical modeling.
- SML is usually applied to classification and regression problems
- Goal: Predict a variable that, for at least a part of our data, is known

Statistical Modeling and Prediction

- Example: media.csv
 - How many days per week respondents turn to different media types (radio, newspaper, tv and Internet) in order to follow the news.
 - Age (in years), gender (coded as female = 0, male = 1), and education (on a 5-point scale)
 - How far do the sociodemographic characteristics of the respondents explain their media use?
- **Statistics:** Typical approach would be to run a regression analysis.
- **OLS:**

$$y = \beta_0 + \beta_1 \cdot \text{age} + \beta_2 \cdot \text{gender} + \beta_3 \cdot \text{education} + \epsilon$$

where y is the number of days per week a respondent uses the Internet to follow the news.

- Also look at R^2 to see how well the model fits the data.

Statistical Modeling and Prediction

- 1 Fit a model using known (labeled) data (statsmodels package).
- 2 Use the fitted model to predict outcomes for new, unseen data.
- 3 Report and interpret: coefficients, standard errors, confidence intervals, p-values and R^2 .

```
df = read.csv("https://cssbook.net/d/media.csv")
mod = lm(formula = "newspaper ~ age + gender",
          data = df)
# summary(mod) would give a lot more info,
# but we only care about the coefficients:
mod
```

Call:

```
lm(formula = "newspaper ~ age + gender", data = df)
```

Coefficients:

(Intercept)	age	gender
-0.08956	0.06762	0.17666

Statistical Modeling and Prediction

- We can now plug in values to make predictions.

$$\hat{y} = -0.0896 + 0.0676 \cdot \text{age} + 0.1767 \cdot \text{gender}$$

```
gender = c(1,0)
age = c(20,40)
newdata = data.frame(age, gender)
predict(mod, newdata)
```

```
      1      2
1.439508 2.615248
```

Note: Prediction model can take the form of *any* function, as long as it takes some characteristics (or “features”) of the cases (in this case, people) as input and returns a prediction.

Model Limitations and Interpretation

- Predictions may be implausible (e.g., negative or >7 days/week).
- Linear models assume:
 - Linearity
 - Independence of errors
 - Homoskedasticity
- These assumptions may not hold in practice.
- Many tasks are better suited for **classification** than regression.
 - Depends on the goal of the analysis.
 - Sometimes predict outcomes, not necessarily an exact value.

Concepts and Principles of SML

- Learn from labeled data to predict outcomes for unseen data.
 - We did a simple version of this with OLS regression.
 - When do we need to use a SML approach?
- Two preconditions:
 - Large dataset
 - Random subset of data with known outcomes (labels)

Machine Learning vs. Statistical Terminology

Machine Learning	Statistics
Feature	Independent variable
Label	Dependent variable
Labeled dataset	Data with known outcomes
Train / fit model	Estimate model
Classifier	Model predicting nominal outcomes

Typical Supervised Learning Workflow

- ① Annotate or label a subset of the data.
 - Should be *class balanced*.
- ② Split the labelled data into **training** and **test** sets.
 - Common split ratios: 50:50 to 80:20.
- ③ Train model on training data.
- ④ Evaluate performance on unseen test data.

Typical Supervised Learning Workflow

Split the labelled data into training and test sets:

```
df = read.csv("https://cssbook.net/d/media.csv")
df = na.omit(df %>% mutate(
  usesinternet=recode(internet,
    .default="user", `0`="non-user"))))

set.seed(42)
df$usesinternet = as.factor(df$usesinternet)
print("How many people used online news at all?")
```

```
[1] "How many people used online news at all?"
```

```
print(table(df$usesinternet))
```

non-user	user
803	1262

Typical Supervised Learning Workflow

Split the labelled data into training and test sets:

```
split = initial_split(df, prop = .8)
traindata = training(split)
testdata  = testing(split)

X_train = select(traindata,
                  c("age", "gender", "education"))
y_train = traindata$usesinternet
X_test  = select(testdata,
                  c("age", "gender", "education"))
y_test  = testdata$usesinternet

glue("We have {nrow(X_train)} training and {nrow(X_test)} test cases.")
```

We have 1652 training and 413 test cases.

Typical Supervised Learning Workflow

We now can train our classifier!


We estimate our model using the training dataset objects `X_train` and `y_train`:

```
myclassifier = train(x = X_train, y = y_train,  
                    method = "naive_bayes")  
y_pred = predict(myclassifier, newdata = X_test)
```

BUT before we can use this classifier, we need to test how capable it is to predict the correct labels, given a set of features.

Evaluating Model Performance

- For this, we need to evaluate the model on the unseen test set.
- We could use the same input data, but this is not strict enough.
- We don't want a model that is only good at predicting its own training data (overfitting).
- We want a model that generalizes well to unseen data.



Evaluating Model Performance

- Use a **confusion matrix** to compare predictions and true labels.

		Predicted Class	
		<i>Positive</i>	<i>Negative</i>
Real / Actual Class	<i>Positive</i>	<div>TP True Positive</div>	<div>FN False Negative</div>
	<i>Negative</i>	<div>FP False Positive</div>	<div>TN True Negative</div>

Precision and Recall Metrics

$$\text{Precision} = \frac{TP}{TP + FP} \quad \text{Recall} = \frac{TP}{TP + FN}$$

- **Precision:** correctness of positive predictions.
- **Recall:** completeness of positive predictions.
- Often, increasing one decreases the other (trade-off).